

---

# Data Model Design for MongoDB

*Release 3.0.4*

## MongoDB Documentation Project

June 19, 2015

## Contents

<b>1 Data Modeling Introduction</b>	<b>3</b>
1.1 Document Structure . . . . .	3
References . . . . .	3
Embedded Data . . . . .	4
1.2 Atomicity of Write Operations . . . . .	4
1.3 Document Growth . . . . .	4
1.4 Data Use and Performance . . . . .	5
1.5 Additional Resources . . . . .	5
<b>2 Data Modeling Concepts</b>	<b>5</b>
2.1 Data Model Design . . . . .	5
Embedded Data Models . . . . .	5
Normalized Data Models . . . . .	6
Additional Resources . . . . .	7
2.2 Operational Factors and Data Models . . . . .	7
Document Growth . . . . .	7
Atomicity . . . . .	8
Sharding . . . . .	8
Indexes . . . . .	8
Large Number of Collections . . . . .	9
Data Lifecycle Management . . . . .	9
2.3 GridFS . . . . .	10
Implement GridFS . . . . .	10
GridFS Collections . . . . .	10
GridFS Index . . . . .	11
Additional Resources . . . . .	11
<b>3 Data Model Examples and Patterns</b>	<b>11</b>
3.1 Model Relationships Between Documents . . . . .	12
Model One-to-One Relationships with Embedded Documents . . . . .	12
Model One-to-Many Relationships with Embedded Documents . . . . .	13
Model One-to-Many Relationships with Document References . . . . .	14
3.2 Model Tree Structures . . . . .	16
Model Tree Structures with Parent References . . . . .	17
Model Tree Structures with Child References . . . . .	18
Model Tree Structures with an Array of Ancestors . . . . .	19

Model Tree Structures with Materialized Paths . . . . .	21
Model Tree Structures with Nested Sets . . . . .	22
3.3 Model Specific Application Contexts . . . . .	23
Model Data for Atomic Operations . . . . .	24
Model Data to Support Keyword Search . . . . .	24
Model Monetary Data . . . . .	26
Model Time Data . . . . .	27
<b>4 Data Model Reference</b> . . . . .	<b>28</b>
4.1 Documents . . . . .	28
Document Format . . . . .	28
Document Structure . . . . .	28
Field Names . . . . .	29
Field Value Limit . . . . .	29
Document Limitations . . . . .	30
The <code>_id</code> Field . . . . .	30
Dot Notation . . . . .	31
Additional Resources . . . . .	31
4.2 Database References . . . . .	31
Manual References . . . . .	32
DBRefs . . . . .	32
4.3 GridFS Reference . . . . .	34
The <code>chunks</code> Collection . . . . .	34
The <code>files</code> Collection . . . . .	35
4.4 ObjectId . . . . .	36
Overview . . . . .	36
ObjectId() . . . . .	36
Examples . . . . .	37
4.5 BSON Types . . . . .	38
Comparison/Sort Order . . . . .	38
ObjectId . . . . .	39
String . . . . .	39
Timestamps . . . . .	39
Date . . . . .	40

---

Data in MongoDB has a *flexible schema*. Collections do not enforce *document* structure. This flexibility gives you data-modeling choices to match your application and its performance requirements.

**[Data Modeling Introduction](#)** (page 3) An introduction to data modeling in MongoDB.

**[Data Modeling Concepts](#)** (page 5) The core documentation detailing the decisions you must make when determining a data model, and discussing considerations that should be taken into account.

**[Data Model Examples and Patterns](#)** (page 11) Examples of possible data models that you can use to structure your MongoDB documents.

**[Data Model Reference](#)** (page 28) Reference material for data modeling for developers of MongoDB applications.

# 1 Data Modeling Introduction

Data in MongoDB has a *flexible schema*. Unlike SQL databases, where you must determine and declare a table's schema before inserting data, MongoDB's *collections* do not enforce *document* structure. This flexibility facilitates the mapping of documents to an entity or an object. Each document can match the data fields of the represented entity, even if the data has substantial variation. In practice, however, the documents in a collection share a similar structure.


The key challenge in data modeling is balancing the needs of the application, the performance characteristics of the database engine, and the data retrieval patterns. When designing data models, always consider the application usage of the data (i.e. queries, updates, and processing of the data) as well as the inherent structure of the data itself.

## 1.1 Document Structure

The key decision in designing data models for MongoDB applications revolves around the structure of documents and how the application represents relationships between data. There are two tools that allow applications to represent these relationships: *references* and *embedded documents*.

### References

References store the relationships between data by including links or *references* from one document to another. Applications can resolve these *references* (page 31) to access the related data. Broadly, these are *normalized* data models.



See *Normalized Data Models* (page 6) for the strengths and weaknesses of using references.

## Embedded Data

Embedded documents capture relationships between data by storing related data in a single document structure. MongoDB documents make it possible to embed document structures in a field or array within a document. These *denormalized* data models allow applications to retrieve and manipulate related data in a single database operation.

```
{  
  _id: <ObjectId1>,  
  username: "123xyz",  
  contact: {  
    phone: "123-456-7890",  
    email: "xyz@example.com"  
  },  
  access: {  
    level: 5,  
    group: "dev"  
  }  
}
```



Embedded sub-document  
Embedded sub-document

See [Embedded Data Models](#) (page 5) for the strengths and weaknesses of embedding documents.

## 1.2 Atomicity of Write Operations

In MongoDB, write operations are atomic at the *document* level, and no single write operation can atomically affect more than one document or more than one collection. A denormalized data model with embedded data combines all related data for a represented entity in a single document. This facilitates atomic write operations since a single write operation can insert or update the data for an entity. Normalizing the data would split the data across multiple collections and would require multiple write operations that are not atomic collectively.

However, schemas that facilitate atomic writes may limit ways that applications can use the data or may limit ways to modify applications. The [Atomicity Considerations](#) (page 8) documentation describes the challenge of designing a schema that balances flexibility and atomicity.

## 1.3 Document Growth

Some updates, such as pushing elements to an array or adding new fields, increase a *document*'s size.

For the MMAPv1 storage engine, if the document size exceeds the allocated space for that document, MongoDB relocates the document on disk. When using the MMAPv1 storage engine, growth consideration can affect the decision to normalize or denormalize data. See [Document Growth Considerations](#) (page 7) for more about planning for and managing document growth for MMAPv1.

## 1.4 Data Use and Performance

When designing a data model, consider how applications will use your database. For instance, if your application only uses recently inserted documents, consider using <http://docs.mongodb.org/manual/core/capped-collections>. Or if your application needs are mainly read operations to a collection, adding indexes to support common queries can improve performance.

See [Operational Factors and Data Models](#) (page 7) for more information on these and other operational considerations that affect data model designs.

## 1.5 Additional Resources

- Thinking in Documents Part 1 (Blog Post)<sup>1</sup>

# 2 Data Modeling Concepts

Consider the following aspects of data modeling in MongoDB:

[Data Model Design](#) (page 5) Presents the different strategies that you can choose from when determining your data model, their strengths and their weaknesses.

[Operational Factors and Data Models](#) (page 7) Details features you should keep in mind when designing your data model, such as lifecycle management, indexing, horizontal scalability, and document growth.

[GridFS](#) (page 10) GridFS is a specification for storing documents that exceeds the *BSON*-document size limit of 16MB.

For a general introduction to data modeling in MongoDB, see the [Data Modeling Introduction](#) (page 3). For example data models, see [Data Modeling Examples and Patterns](#) (page 11).

## 2.1 Data Model Design

Effective data models support your application needs. The key consideration for the structure of your documents is the decision to [embed](#) (page 5) or to [use references](#) (page 6).

### Embedded Data Models

With MongoDB, you may embed related data in a single structure or document. These schema are generally known as “denormalized” models, and take advantage of MongoDB’s rich documents. Consider the following diagram:

Embedded data models allow applications to store related pieces of information in the same database record. As a result, applications may need to issue fewer queries and updates to complete common operations.

In general, use embedded data models when:

- you have “contains” relationships between entities. See [Model One-to-One Relationships with Embedded Documents](#) (page 12).
- you have one-to-many relationships between entities. In these relationships the “many” or child documents always appear with or are viewed in the context of the “one” or parent documents. See [Model One-to-Many Relationships with Embedded Documents](#) (page 13).

---

<sup>1</sup><https://www.mongodb.com/blog/post/thinking-documents-part-1?jmp=docs>

```
{
  _id: <ObjectId1>,
  username: "123xyz",
  contact: {
    phone: "123-456-7890",
    email: "xyz@example.com"
  },
  access: {
    level: 5,
    group: "dev"
  }
}
```

In general, embedding provides better performance for read operations, as well as the ability to request and retrieve related data in a single database operation. Embedded data models make it possible to update related data in a single atomic write operation.

However, embedding related data in documents may lead to situations where documents grow after creation. With the MMAPv1 storage engine, document growth can impact write performance and lead to data fragmentation.

In version 3.0.0, MongoDB uses *power-of-2-allocation* as the default allocation strategy for MMAPv1 in order to account for document growth, minimizing the likelihood of data fragmentation. See *power-of-2-allocation* for details. Furthermore, documents in MongoDB must be smaller than the maximum BSON document size. For bulk binary data, consider [GridFS](#) (page 10).

To interact with embedded documents, use *dot notation* to “reach into” embedded documents. See *query for data in arrays* and *query data in embedded documents* for more examples on accessing data in arrays and embedded documents.

## Normalized Data Models


Normalized data models describe relationships using [references](#) (page 31) between documents.

In general, use normalized data models:

- when embedding would result in duplication of data but would not provide sufficient read performance advantages to outweigh the implications of the duplication.
- to represent more complex many-to-many relationships.
- to model large hierarchical data sets.

References provides more flexibility than embedding. However, client-side applications must issue follow-up queries to resolve the references. In other words, normalized data models can require more round trips to the server.

See [Model One-to-Many Relationships with Document References](#) (page 14) for an example of referencing. For examples of various tree models using references, see [Model Tree Structures](#) (page 16).



## Additional Resources

- Thinking in Documents Part 1 (Blog Post)<sup>2</sup>
- Thinking in Documents (Presentation)<sup>3</sup>
- Schema Design for Time Series Data (Presentation)<sup>4</sup>
- Socialite, the Open Source Status Feed - Storing a Social Graph (Presentation)<sup>5</sup>
- MongoDB Rapid Start Consultation Services<sup>6</sup>

## 2.2 Operational Factors and Data Models

Modeling application data for MongoDB depends on both the data itself, as well as the characteristics of MongoDB itself. For example, different data models may allow applications to use more efficient queries, increase the throughput of insert and update operations, or distribute activity to a sharded cluster more effectively.

These factors are *operational* or address requirements that arise outside of the application but impact the performance of MongoDB based applications. When developing a data model, analyze all of your application's read operations and write operations in conjunction with the following considerations.

### Document Growth

Changed in version 3.0.0.

<sup>2</sup><https://www.mongodb.com/blog/post/thinking-documents-part-1?jmp=docs>

<sup>3</sup><http://www.mongodb.com/presentations/webinar-back-basics-1-thinking-documents?jmp=docs>

<sup>4</sup><http://www.mongodb.com/presentations/webinar-time-series-data-mongodb?jmp=docs>

<sup>5</sup><http://www.mongodb.com/presentations/socialite-open-source-status-feed-part-2-managing-social-graph?jmp=docs>

<sup>6</sup>[https://www.mongodb.com/products/consulting?jmp=docs#rapid\\_start](https://www.mongodb.com/products/consulting?jmp=docs#rapid_start)

Some updates to documents can increase the size of documents. These updates include pushing elements to an array (i.e. `$push`) and adding new fields to a document.

When using the MMAPv1 storage engine, document growth can be a consideration for your data model. For MMAPv1, if the document size exceeds the allocated space for that document, MongoDB will relocate the document on disk. With MongoDB 3.0.0, however, the default use of the *power-of-2-allocation* minimizes the occurrences of such re-allocations as well as allows for the effective reuse of the freed record space.

When using MMAPv1, if your applications require updates that will frequently cause document growth to exceed the current power of 2 allocation, you may want to refactor your data model to use references between data in distinct documents rather than a denormalized data model.

You may also use a *pre-allocation* strategy to explicitly avoid document growth. Refer to the [Pre-Aggregated Reports Use Case<sup>7</sup>](#) for an example of the *pre-allocation* approach to handling document growth.

See <http://docs.mongodb.org/manual/core/storage> for more information on MongoDB's storage model and record allocation strategies.

## Atomicity

In MongoDB, operations are atomic at the *document* level. No **single** write operation can change more than one document. Operations that modify more than a single document in a collection still operate on one document at a time.

<sup>8</sup> Ensure that your application stores all fields with atomic dependency requirements in the same document. If the application can tolerate non-atomic updates for two pieces of data, you can store these data in separate documents.

A data model that embeds related data in a single document facilitates these kinds of atomic operations. For data models that store references between related pieces of data, the application must issue separate read and write operations to retrieve and modify these related pieces of data.

See [Model Data for Atomic Operations](#) (page 24) for an example data model that provides atomic updates for a single document.

## Sharding

MongoDB uses *sharding* to provide horizontal scaling. These clusters support deployments with large data sets and high-throughput operations. Sharding allows users to *partition* a *collection* within a database to distribute the collection's documents across a number of `mongod` instances or *shards*.

To distribute data and application traffic in a sharded collection, MongoDB uses the *shard key*. Selecting the proper *shard key* has significant implications for performance, and can enable or prevent query isolation and increased write capacity. It is important to consider carefully the field or fields to use as the shard key.

See <http://docs.mongodb.org/manual/core/sharding-introduction> and <http://docs.mongodb.org/manual/core/sharding-shard-key> for more information.

## Indexes

Use indexes to improve performance for common queries. Build indexes on fields that appear often in queries and for all operations that return sorted results. MongoDB automatically creates a unique index on the `_id` field.

As you create indexes, consider the following behaviors of indexes:

- Each index requires at least 8KB of data space.

<sup>7</sup><http://docs.mongodb.org/ecosystem/use-cases/pre-aggregated-reports>

<sup>8</sup> Document-level atomic operations include all operations within a single MongoDB document record: operations that affect multiple embedded documents within that single record are still atomic.

- Adding an index has some negative performance impact for write operations. For collections with high write-to-read ratio, indexes are expensive since each insert must also update any indexes.
- Collections with high read-to-write ratio often benefit from additional indexes. Indexes do not affect un-indexed read operations.
- When active, each index consumes disk space and memory. This usage can be significant and should be tracked for capacity planning, especially for concerns over working set size.

See <http://docs.mongodb.org/manual/applications/indexes> for more information on indexes as well as <http://docs.mongodb.org/manual/tutorial/analyze-query-plan/>. Additionally, the MongoDB database profiler may help identify inefficient queries.

## Large Number of Collections

In certain situations, you might choose to store related information in several collections rather than in a single collection.

Consider a sample collection `logs` that stores log documents for various environment and applications. The `logs` collection contains documents of the following form:

```
{ log: "dev", ts: ..., info: ... }
{ log: "debug", ts: ..., info: ... }
```

If the total number of documents is low, you may group documents into collection by type. For `logs`, consider maintaining distinct log collections, such as `logs_dev` and `logs_debug`. The `logs_dev` collection would contain only the documents related to the dev environment.

Generally, having a large number of collections has no significant performance penalty and results in very good performance. Distinct collections are very important for high-throughput batch processing.

When using models that have a large number of collections, consider the following behaviors:

- Each collection has a certain minimum overhead of a few kilobytes.
- Each index, including the index on `_id`, requires at least 8KB of data space.
- For each *database*, a single namespace file (i.e. `<database>.ns`) stores all meta-data for that database, and each index and collection has its own entry in the namespace file. MongoDB places limits on the size of namespace files.
- MongoDB using the mmapv1 storage engine has limits on the number of namespaces. You may wish to know the current number of namespaces in order to determine how many additional namespaces the database can support. To get the current number of namespaces, run the following in the `mongo` shell:

```
db.system.namespaces.count()
```

The limit on the number of namespaces depend on the `<database>.ns` size. The namespace file defaults to 16 MB.

To change the size of the *new* namespace file, start the server with the option `--nssize <new size MB>`. For existing databases, after starting up the server with `--nssize`, run the `db.repairDatabase()` command from the `mongo` shell. For impacts and considerations on running `db.repairDatabase()`, see `repairDatabase`.

## Data Lifecycle Management

Data modeling decisions should take data lifecycle management into consideration.

The Time to Live or TTL feature of collections expires documents after a period of time. Consider using the TTL feature if your application requires some data to persist in the database for a limited period of time.

Additionally, if your application only uses recently inserted documents, consider <http://docs.mongodb.org/manual/core/capped-collections>. Capped collections provide first-in-first-out (FIFO) management of inserted documents and efficiently support operations that insert and read documents based on insertion order.

## 2.3 GridFS

*GridFS* is a specification for storing and retrieving files that exceed the  *BSON-document size limit of 16MB*.

Instead of storing a file in a single document, GridFS divides a file into parts, or chunks,<sup>9</sup> and stores each of those chunks as a separate document. By default GridFS limits chunk size to 255k. GridFS uses two collections to store files. One collection stores the file chunks, and the other stores file metadata.

When you query a GridFS store for a file, the driver or client will reassemble the chunks as needed. You can perform range queries on files stored through GridFS. You also can access information from arbitrary sections of files, which allows you to “skip” into the middle of a video or audio file.

GridFS is useful not only for storing files that exceed 16MB but also for storing any files for which you want access without having to load the entire file into memory. For more information on the indications of GridFS, see *faq-developers-when-to-use-gridfs*.

Changed in version 2.4.10: The default chunk size changed from 256k to 255k.

### Implement GridFS

To store and retrieve files using *GridFS*, use either of the following:

- A MongoDB driver. See the `drivers` documentation for information on using GridFS with your driver.
- The `mongofiles` command-line tool in the `mongo` shell. See the `mongofiles` reference for complete documentation.

### GridFS Collections

*GridFS* stores files in two collections:

- `chunks` stores the binary chunks. For details, see *The chunks Collection* (page 34).
- `files` stores the file’s metadata. For details, see *The files Collection* (page 35).

GridFS places the collections in a common bucket by prefixing each with the bucket name. By default, GridFS uses two collections with names prefixed by `fs` bucket:

- `fs.files`
- `fs.chunks`

You can choose a different bucket name than `fs`, and create multiple buckets in a single database.

Each document in the `chunks` collection represents a distinct chunk of a file as represented in the GridFS store. Each chunk is identified by its unique `ObjectId` stored in its `_id` field.

For descriptions of all fields in the `chunks` and `files` collections, see *GridFS Reference* (page 34).

<sup>9</sup> The use of the term `chunks` in the context of GridFS is not related to the use of the term `chunks` in the context of sharding.

## GridFS Index

*GridFS* uses a *unique, compound* index on the `chunks` collection for the `files_id` and `n` fields. The `files_id` field contains the `_id` of the chunk's "parent" document. The `n` field contains the sequence number of the chunk. GridFS numbers all chunks, starting with 0. For descriptions of the documents and fields in the `chunks` collection, see [GridFS Reference](#) (page 34).

The GridFS index allows efficient retrieval of chunks using the `files_id` and `n` values, as shown in the following example:

```
cursor = db.fs.chunks.find({files_id: myFileID}).sort({n:1});
```

See the relevant `driver` documentation for the specific behavior of your GridFS application. If your driver does not create this index, issue the following operation using the `mongo` shell:

```
db.fs.chunks.createIndex( { files_id: 1, n: 1 }, { unique: true } );
```

## Additional Resources

- Building MongoDB Applications with Binary Files Using GridFS: Part 1<sup>10</sup>
- Building MongoDB Applications with Binary Files Using GridFS: Part 2<sup>11</sup>

## 3 Data Model Examples and Patterns

The following documents provide overviews of various data modeling patterns and common schema design considerations:

[Model Relationships Between Documents](#) (page 12) Examples for modeling relationships between documents.

[Model One-to-One Relationships with Embedded Documents](#) (page 12) Presents a data model that uses *embedded documents* (page 5) to describe one-to-one relationships between connected data.

[Model One-to-Many Relationships with Embedded Documents](#) (page 13) Presents a data model that uses *embedded documents* (page 5) to describe one-to-many relationships between connected data.

[Model One-to-Many Relationships with Document References](#) (page 14) Presents a data model that uses *references* (page 6) to describe one-to-many relationships between documents.

[Model Tree Structures](#) (page 16) Examples for modeling tree structures.

[Model Tree Structures with Parent References](#) (page 17) Presents a data model that organizes documents in a tree-like structure by storing *references* (page 6) to "parent" nodes in "child" nodes.

[Model Tree Structures with Child References](#) (page 18) Presents a data model that organizes documents in a tree-like structure by storing *references* (page 6) to "child" nodes in "parent" nodes.

See [Model Tree Structures](#) (page 16) for additional examples of data models for tree structures.

[Model Specific Application Contexts](#) (page 23) Examples for models for specific application contexts.

[Model Data for Atomic Operations](#) (page 24) Illustrates how embedding fields related to an atomic update within the same document ensures that the fields are in sync.

[Model Data to Support Keyword Search](#) (page 24) Describes one method for supporting keyword search by storing keywords in an array in the same document as the text field. Combined with a multi-key index, this pattern can support application's keyword search operations.

<sup>10</sup><http://www.mongodb.com/blog/post/building-mongodb-applications-binary-files-using-gridfs-part-1?jmp=docs>

<sup>11</sup><http://www.mongodb.com/blog/post/building-mongodb-applications-binary-files-using-gridfs-part-2?jmp=docs>

## 3.1 Model Relationships Between Documents

**Model One-to-One Relationships with Embedded Documents (page 12)** Presents a data model that uses [embedded documents](#) (page 5) to describe one-to-one relationships between connected data.

**Model One-to-Many Relationships with Embedded Documents (page 13)** Presents a data model that uses [embedded documents](#) (page 5) to describe one-to-many relationships between connected data.

**Model One-to-Many Relationships with Document References (page 14)** Presents a data model that uses [references](#) (page 6) to describe one-to-many relationships between documents.

### Model One-to-One Relationships with Embedded Documents

#### Overview

Data in MongoDB has a *flexible schema*. Collections do not enforce *document* structure. Decisions that affect how you model data can affect application performance and database capacity. See [Data Modeling Concepts](#) (page 5) for a full high level overview of data modeling in MongoDB.

This document describes a data model that uses [embedded](#) (page 5) documents to describe relationships between connected data.

#### Pattern

Consider the following example that maps patron and address relationships. The example illustrates the advantage of embedding over referencing if you need to view one data entity in context of the other. In this one-to-one relationship between patron and address data, the address belongs to the patron.

In the normalized data model, the address document contains a reference to the patron document.

```
{  
  _id: "joe",  
  name: "Joe Bookreader"  
}  
  
{  
  patron_id: "joe",  
  street: "123 Fake Street",  
  city: "Faketown",  
  state: "MA",  
  zip: "12345"  
}
```

If the address data is frequently retrieved with the name information, then with referencing, your application needs to issue multiple queries to resolve the reference. The better data model would be to embed the address data in the patron data, as in the following document:

```
{  
  _id: "joe",  
  name: "Joe Bookreader",  
  address: {  
    street: "123 Fake Street",  
    city: "Faketown",  
    state: "MA",  
    zip: "12345"  
  }  
}
```

With the embedded data model, your application can retrieve the complete patron information with one query.

## Model One-to-Many Relationships with Embedded Documents

### Overview

Data in MongoDB has a *flexible schema*. Collections do not enforce *document* structure. Decisions that affect how you model data can affect application performance and database capacity. See [Data Modeling Concepts](#) (page 5) for a full high level overview of data modeling in MongoDB.

This document describes a data model that uses [embedded](#) (page 5) documents to describe relationships between connected data.

### Pattern

Consider the following example that maps patron and multiple address relationships. The example illustrates the advantage of embedding over referencing if you need to view many data entities in context of another. In this one-to-many relationship between patron and address data, the patron has multiple address entities.

In the normalized data model, the address documents contain a reference to the patron document.

```
{  
  _id: "joe",  
  name: "Joe Bookreader"  
}  
  
{  
  patron_id: "joe",  
  street: "123 Fake Street",  
  city: "Faketown",  
  state: "MA",  
  zip: "12345"  
}  
  
{  
  patron_id: "joe",  
  street: "1 Some Other Street",  
  city: "Boston",  
  state: "MA",  
  zip: "12345"  
}
```

If your application frequently retrieves the address data with the name information, then your application needs to issue multiple queries to resolve the references. A more optimal schema would be to embed the address data entities in the patron data, as in the following document:

```
{  
  _id: "joe",  
  name: "Joe Bookreader",  
  addresses: [  
    {  
      street: "123 Fake Street",  
      city: "Faketown",  
      state: "MA",  
      zip: "12345"  
    },  
  ]
```

```

        {
            street: "1 Some Other Street",
            city: "Boston",
            state: "MA",
            zip: "12345"
        }
    ]
}

```

With the embedded data model, your application can retrieve the complete patron information with one query.

## Model One-to-Many Relationships with Document References

### Overview

Data in MongoDB has a *flexible schema*. *Collections* do not enforce *document* structure. Decisions that affect how you model data can affect application performance and database capacity. See [Data Modeling Concepts](#) (page 5) for a full high level overview of data modeling in MongoDB.

This document describes a data model that uses [references](#) (page 6) between documents to describe relationships between connected data.

### Pattern

Consider the following example that maps publisher and book relationships. The example illustrates the advantage of referencing over embedding to avoid repetition of the publisher information.

Embedding the publisher document inside the book document would lead to **repetition** of the publisher data, as the following documents show:

```

{
    title: "MongoDB: The Definitive Guide",
    author: [ "Kristina Chodorow", "Mike Dirolf" ],
    published_date: ISODate("2010-09-24"),
    pages: 216,
    language: "English",
    publisher: {
        name: "O'Reilly Media",
        founded: 1980,
        location: "CA"
    }
}

{
    title: "50 Tips and Tricks for MongoDB Developer",
    author: "Kristina Chodorow",
    published_date: ISODate("2011-05-06"),
    pages: 68,
    language: "English",
    publisher: {
        name: "O'Reilly Media",
        founded: 1980,
        location: "CA"
    }
}

```

To avoid repetition of the publisher data, use *references* and keep the publisher information in a separate collection from the book collection.

When using references, the growth of the relationships determine where to store the reference. If the number of books per publisher is small with limited growth, storing the book reference inside the publisher document may sometimes be useful. Otherwise, if the number of books per publisher is unbounded, this data model would lead to mutable, growing arrays, as in the following example:

```
{  
    name: "O'Reilly Media",  
    founded: 1980,  
    location: "CA",  
    books: [123456789, 234567890, ...]  
}  
  
{  
    _id: 123456789,  
    title: "MongoDB: The Definitive Guide",  
    author: [ "Kristina Chodorow", "Mike Dirolf" ],  
    published_date: ISODate("2010-09-24"),  
    pages: 216,  
    language: "English"  
}  
  
{  
    _id: 234567890,  
    title: "50 Tips and Tricks for MongoDB Developer",  
    author: "Kristina Chodorow",  
    published_date: ISODate("2011-05-06"),  
    pages: 68,  
    language: "English"  
}
```


To avoid mutable, growing arrays, store the publisher reference inside the book document:

```
{  
    _id: "oreilly",  
    name: "O'Reilly Media",  
    founded: 1980,  
    location: "CA"  
}  
  
{  
    _id: 123456789,  
    title: "MongoDB: The Definitive Guide",  
    author: [ "Kristina Chodorow", "Mike Dirolf" ],  
    published_date: ISODate("2010-09-24"),  
    pages: 216,  
    language: "English",  
    publisher_id: "oreilly"  
}  
  
{  
    _id: 234567890,  
    title: "50 Tips and Tricks for MongoDB Developer",  
    author: "Kristina Chodorow",  
    published_date: ISODate("2011-05-06"),  
    pages: 68,  
    language: "English",  
}
```

```
    publisher_id: "oreilly"  
}
```

## 3.2 Model Tree Structures

MongoDB allows various ways to use tree data structures to model large hierarchical or nested data relationships.



***Model Tree Structures with Parent References*** (page 17) Presents a data model that organizes documents in a tree-like structure by storing *references* (page 6) to “parent” nodes in “child” nodes.

***Model Tree Structures with Child References*** (page 18) Presents a data model that organizes documents in a tree-like structure by storing *references* (page 6) to “child” nodes in “parent” nodes.

***Model Tree Structures with an Array of Ancestors*** (page 19) Presents a data model that organizes documents in a tree-like structure by storing *references* (page 6) to “parent” nodes and an array that stores all ancestors.

***Model Tree Structures with Materialized Paths*** (page 21) Presents a data model that organizes documents in a tree-like structure by storing full relationship paths between documents. In addition to the tree node, each document stores the `_id` of the nodes ancestors or path as a string.

***Model Tree Structures with Nested Sets*** (page 22) Presents a data model that organizes documents in a tree-like structure using the *Nested Sets* pattern. This optimizes discovering subtrees at the expense of tree mutability.

## Model Tree Structures with Parent References

### Overview


Data in MongoDB has a *flexible schema*. Collections do not enforce *document* structure. Decisions that affect how you model data can affect application performance and database capacity. See [Data Modeling Concepts](#) (page 5) for a full high level overview of data modeling in MongoDB.

This document describes a data model that describes a tree-like structure in MongoDB documents by storing *references* (page 6) to “parent” nodes in children nodes.

### Pattern

The *Parent References* pattern stores each tree node in a document; in addition to the tree node, the document stores the id of the node’s parent.

Consider the following hierarchy of categories:



The following example models the tree using *Parent References*, storing the reference to the parent category in the field `parent`:

```
db.categories.insert( { _id: "MongoDB", parent: "Databases" } )
db.categories.insert( { _id: "dbm", parent: "Databases" } )
db.categories.insert( { _id: "Databases", parent: "Programming" } )
db.categories.insert( { _id: "Languages", parent: "Programming" } )
```

```
db.categories.insert( { _id: "Programming", parent: "Books" } )
db.categories.insert( { _id: "Books", parent: null } )
```

- The query to retrieve the parent of a node is fast and straightforward:

```
db.categories.findOne( { _id: "MongoDB" } ).parent
```

- You can create an index on the field `parent` to enable fast search by the parent node:

```
db.categories.createIndex( { parent: 1 } )
```

- You can query by the `parent` field to find its immediate children nodes:

```
db.categories.find( { parent: "Databases" } )
```

The *Parent Links* pattern provides a simple solution to tree storage but requires multiple queries to retrieve subtrees.

## Model Tree Structures with Child References

### Overview

Data in MongoDB has a *flexible schema*. *Collections* do not enforce *document* structure. Decisions that affect how you model data can affect application performance and database capacity. See [Data Modeling Concepts](#) (page 5) for a full high level overview of data modeling in MongoDB.

This document describes a data model that describes a tree-like structure in MongoDB documents by storing *references* (page 6) in the parent-nodes to children nodes.

### Pattern

The *Child References* pattern stores each tree node in a document; in addition to the tree node, document stores in an array the id(s) of the node's children.

Consider the following hierarchy of categories:

The following example models the tree using *Child References*, storing the reference to the node's children in the field `children`:

```
db.categories.insert( { _id: "MongoDB", children: [] } )
db.categories.insert( { _id: "dbm", children: [] } )
db.categories.insert( { _id: "Databases", children: [ "MongoDB", "dbm" ] } )
db.categories.insert( { _id: "Languages", children: [] } )
db.categories.insert( { _id: "Programming", children: [ "Databases", "Languages" ] } )
db.categories.insert( { _id: "Books", children: [ "Programming" ] } )
```

- The query to retrieve the immediate children of a node is fast and straightforward:


```
db.categories.findOne( { _id: "Databases" } ).children
```

- You can create an index on the field `children` to enable fast search by the child nodes:

```
db.categories.createIndex( { children: 1 } )
```

- You can query for a node in the `children` field to find its parent node as well as its siblings:

```
db.categories.find( { children: "MongoDB" } )
```



The *Child References* pattern provides a suitable solution to tree storage as long as no operations on subtrees are necessary. This pattern may also provide a suitable solution for storing graphs where a node may have multiple parents.

## Model Tree Structures with an Array of Ancestors

### Overview

Data in MongoDB has a *flexible schema*. *Collections* do not enforce *document* structure. Decisions that affect how you model data can affect application performance and database capacity. See *Data Modeling Concepts* (page 5) for a full high level overview of data modeling in MongoDB.


This document describes a data model that describes a tree-like structure in MongoDB documents using *references* (page 6) to parent nodes and an array that stores all ancestors.

### Pattern

The *Array of Ancestors* pattern stores each tree node in a document; in addition to the tree node, document stores in an array the id(s) of the node's ancestors or path.

Consider the following hierarchy of categories:

The following example models the tree using *Array of Ancestors*. In addition to the `ancestors` field, these documents also store the reference to the immediate parent category in the `parent` field:



```

db.categories.insert( { _id: "MongoDB", ancestors: [ "Books", "Programming", "Databases" ], parent: "Databases" } )
db.categories.insert( { _id: "dbm", ancestors: [ "Books", "Programming", "Databases" ], parent: "Databases" } )
db.categories.insert( { _id: "Databases", ancestors: [ "Books", "Programming" ], parent: "Programming" } )
db.categories.insert( { _id: "Languages", ancestors: [ "Books", "Programming" ], parent: "Programming" } )
db.categories.insert( { _id: "Programming", ancestors: [ "Books" ], parent: "Books" } )
db.categories.insert( { _id: "Books", ancestors: [ ], parent: null } )
  
```

- The query to retrieve the ancestors or path of a node is fast and straightforward:

```
db.categories.findOne( { _id: "MongoDB" } ).ancestors
```

- You can create an index on the field `ancestors` to enable fast search by the ancestors nodes:

```
db.categories.createIndex( { ancestors: 1 } )
```

- You can query by the field `ancestors` to find all its descendants:

```
db.categories.find( { ancestors: "Programming" } )
```

The *Array of Ancestors* pattern provides a fast and efficient solution to find the descendants and the ancestors of a node by creating an index on the elements of the `ancestors` field. This makes *Array of Ancestors* a good choice for working with subtrees.

The *Array of Ancestors* pattern is slightly slower than the *Materialized Paths* (page 21) pattern but is more straightforward to use.

## Model Tree Structures with Materialized Paths

### Overview


Data in MongoDB has a *flexible schema*. Collections do not enforce *document* structure. Decisions that affect how you model data can affect application performance and database capacity. See [Data Modeling Concepts](#) (page 5) for a full high level overview of data modeling in MongoDB.

This document describes a data model that describes a tree-like structure in MongoDB documents by storing full relationship paths between documents.

### Pattern

The *Materialized Paths* pattern stores each tree node in a document; in addition to the tree node, document stores as a string the id(s) of the node's ancestors or path. Although the *Materialized Paths* pattern requires additional steps of working with strings and regular expressions, the pattern also provides more flexibility in working with the path, such as finding nodes by partial paths.

Consider the following hierarchy of categories:



The following example models the tree using *Materialized Paths*, storing the path in the field `path`; the path string uses the comma `,` as a delimiter:

```
db.categories.insert( { _id: "Books", path: null } )
db.categories.insert( { _id: "Programming", path: ",Books," } )
db.categories.insert( { _id: "Databases", path: ",Books,Programming," } )
```

```
db.categories.insert( { _id: "Languages", path: ",Books,Programming," } )
db.categories.insert( { _id: "MongoDB", path: ",Books,Programming,Databases," } )
db.categories.insert( { _id: "dbm", path: ",Books,Programming,Databases," } )
```

- You can query to retrieve the whole tree, sorting by the field path:

```
db.categories.find().sort( { path: 1 } )
```

- You can use regular expressions on the path field to find the descendants of Programming:

```
db.categories.find( { path: /,Programming,/ } )
```

- You can also retrieve the descendants of Books where the Books is also at the topmost level of the hierarchy:

```
db.categories.find( { path: /^,Books,/ } )
```

- To create an index on the field path use the following invocation:

```
db.categories.createIndex( { path: 1 } )
```

This index may improve performance depending on the query:

- For queries from the root Books sub-tree (e.g. `http://docs.mongodb.org/manual/^,Books,/` or `http://docs.mongodb.org/manual/^,Books,Programming,/`), an index on the path field improves the query performance significantly.
- For queries of sub-trees where the path from the root is not provided in the query (e.g. `http://docs.mongodb.org/manual/,Databases,/`), or similar queries of sub-trees, where the node might be in the middle of the indexed string, the query must inspect the entire index.

For these queries an index *may* provide some performance improvement *if* the index is significantly smaller than the entire collection.

## Model Tree Structures with Nested Sets

### Overview

Data in MongoDB has a *flexible schema*. Collections do not enforce *document* structure. Decisions that affect how you model data can affect application performance and database capacity. See [Data Modeling Concepts](#) (page 5) for a full high level overview of data modeling in MongoDB.


This document describes a data model that describes a tree like structure that optimizes discovering subtrees at the expense of tree mutability.

### Pattern

The *Nested Sets* pattern identifies each node in the tree as stops in a round-trip traversal of the tree. The application visits each node in the tree twice; first during the initial trip, and second during the return trip. The *Nested Sets* pattern stores each tree node in a document; in addition to the tree node, document stores the id of node's parent, the node's initial stop in the `left` field, and its return stop in the `right` field.

Consider the following hierarchy of categories:

The following example models the tree using *Nested Sets*:



```

db.categories.insert( { _id: "Books", parent: 0, left: 1, right: 12 } )
db.categories.insert( { _id: "Programming", parent: "Books", left: 2, right: 11 } )
db.categories.insert( { _id: "Languages", parent: "Programming", left: 3, right: 4 } )
db.categories.insert( { _id: "Databases", parent: "Programming", left: 5, right: 10 } )
db.categories.insert( { _id: "MongoDB", parent: "Databases", left: 6, right: 7 } )
db.categories.insert( { _id: "dbm", parent: "Databases", left: 8, right: 9 } )
    
```

You can query to retrieve the descendants of a node:

```

var databaseCategory = db.categories.findOne( { _id: "Databases" } );
db.categories.find( { left: { $gt: databaseCategory.left }, right: { $lt: databaseCategory.right } } )
    
```

The *Nested Sets* pattern provides a fast and efficient solution for finding subtrees but is inefficient for modifying the tree structure. As such, this pattern is best for static trees that do not change.

### 3.3 Model Specific Application Contexts

**Model Data for Atomic Operations** (page 24) Illustrates how embedding fields related to an atomic update within the same document ensures that the fields are in sync.

**Model Data to Support Keyword Search** (page 24) Describes one method for supporting keyword search by storing keywords in an array in the same document as the text field. Combined with a multi-key index, this pattern can support application's keyword search operations.

**Model Monetary Data** (page 26) Describes two methods to model monetary data in MongoDB.

**Model Time Data** (page 27) Describes how to deal with local time in MongoDB.

## Model Data for Atomic Operations

### Pattern

In MongoDB, write operations, e.g. `db.collection.update()`, `db.collection.findAndModify()`, `db.collection.remove()`, are atomic on the level of a single document. For fields that must be updated together, embedding the fields within the same document ensures that the fields can be updated atomically.

For example, consider a situation where you need to maintain information on books, including the number of copies available for checkout as well as the current checkout information.

The available copies of the book and the checkout information should be in sync. As such, embedding the `available` field and the `checkout` field within the same document ensures that you can update the two fields atomically.

```
{  
  _id: 123456789,  
  title: "MongoDB: The Definitive Guide",  
  author: [ "Kristina Chodorow", "Mike Dirolf" ],  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English",  
  publisher_id: "oreilly",  
  available: 3,  
  checkout: [ { by: "joe", date: ISODate("2012-10-15") } ]  
}
```

Then to update with new checkout information, you can use the `db.collection.update()` method to atomically update both the `available` field and the `checkout` field:

```
db.books.update (  
  { _id: 123456789, available: { $gt: 0 } },  
  {  
    $inc: { available: -1 },  
    $push: { checkout: { by: "abc", date: new Date() } }  
  }  
)
```

The operation returns a `WriteResult()` object that contains information on the status of the operation:

```
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
```

The `nMatched` field shows that 1 document matched the update condition, and `nModified` shows that the operation updated 1 document.

If no document matched the update condition, then `nMatched` and `nModified` would be 0 and would indicate that you could not check out the book.

## Model Data to Support Keyword Search

---

**Note:** Keyword search is *not* the same as text search or full text search, and does not provide stemming or other text-processing features. See the [Limitations of Keyword Indexes](#) (page 25) section for more information.

In 2.4, MongoDB provides a text search feature. See <http://docs.mongodb.org/manual/core/index-text> for more information.

---

If your application needs to perform queries on the content of a field that holds text you can perform exact matches on the text or use `$regex` to use regular expression pattern matches. However, for many operations on text, these methods do not satisfy application requirements.

This pattern describes one method for supporting keyword search using MongoDB to support application search functionality, that uses keywords stored in an array in the same document as the text field. Combined with a *multi-key index*, this pattern can support application's keyword search operations.

## Pattern

To add structures to your document to support keyword-based queries, create an array field in your documents and add the keywords as strings in the array. You can then create a *multi-key index* on the array and create queries that select values from the array.

---

### Example

Given a collection of library volumes that you want to provide topic-based search. For each volume, you add the array `topics`, and you add as many keywords as needed for a given volume.

For the *Moby-Dick* volume you might have the following document:

```
{ title : "Moby-Dick" ,  
  author : "Herman Melville" ,  
  published : 1851 ,  
  ISBN : 0451526996 ,  
  topics : [ "whaling" , "allegory" , "revenge" , "American" ,  
            "novel" , "nautical" , "voyage" , "Cape Cod" ]  
}
```

You then create a multi-key index on the `topics` array:

```
db.volumes.createIndex( { topics: 1 } )
```

The multi-key index creates separate index entries for each keyword in the `topics` array. For example the index contains one entry for `whaling` and another for `allegory`.

You then query based on the keywords. For example:

```
db.volumes.findOne( { topics : "voyage" } , { title: 1 } )
```

---

**Note:** An array with a large number of elements, such as one with several hundreds or thousands of keywords will incur greater indexing costs on insertion.

---

## Limitations of Keyword Indexes

MongoDB can support keyword searches using specific data models and *multi-key indexes*; however, these keyword indexes are not sufficient or comparable to full-text products in the following respects:

- *Stemming*. Keyword queries in MongoDB can not parse keywords for root or related words.
- *Synonyms*. Keyword-based search features must provide support for synonym or related queries in the application layer.
- *Ranking*. The keyword look ups described in this document do not provide a way to weight results.

- **Asynchronous Indexing.** MongoDB builds indexes synchronously, which means that the indexes used for keyword indexes are always current and can operate in real-time. However, asynchronous bulk indexes may be more efficient for some kinds of content and workloads.

## Model Monetary Data

### Overview

MongoDB stores numeric data as either IEEE 754 standard 64-bit floating point numbers or as 32-bit or 64-bit signed integers. Applications that handle monetary data often require capturing fractional units of currency. However, arithmetic on floating point numbers, as implemented in modern hardware, often does not conform to requirements for monetary arithmetic. In addition, some fractional numeric quantities, such as one third and one tenth, have no exact representation in binary floating point numbers.

---

**Note:** Arithmetic mentioned on this page refers to server-side arithmetic performed by mongod or mongos, and not to client-side arithmetic.

---

This document describes two ways to model monetary data in MongoDB:

- [Exact Precision](#) (page 26) which multiplies the monetary value by a power of 10.
- [Arbitrary Precision](#) (page 27) which uses two fields for the value: one field to store the exact monetary value as a non-numeric and another field to store a floating point approximation of the value.

### Use Cases for Exact Precision Model

If you regularly need to perform server-side arithmetic on monetary data, the exact precision model may be appropriate. For instance:

- If you need to query the database for exact, mathematically valid matches, use [Exact Precision](#) (page 26).
- If you need to be able to do server-side arithmetic, e.g., \$inc, \$mul, and aggregation framework arithmetic, use [Exact Precision](#) (page 26).

### Use Cases for Arbitrary Precision Model

If there is no need to perform server-side arithmetic on monetary data, modeling monetary data using the arbitrary precision model may be suitable. For instance:

- If you need to handle arbitrary or unforeseen number of precision, see [Arbitrary Precision](#) (page 27).
- If server-side approximations are sufficient, possibly with client-side post-processing, see [Arbitrary Precision](#) (page 27).

### Exact Precision

To model monetary data using the exact precision model:

1. Determine the maximum precision needed for the monetary value. For example, your application may require precision down to the tenth of one cent for monetary values in USD currency.
2. Convert the monetary value into an integer by multiplying the value by a power of 10 that ensures the maximum precision needed becomes the least significant digit of the integer. For example, if the required maximum precision is the tenth of one cent, multiply the monetary value by 1000.

3. Store the converted monetary value.

For example, the following scales 9.99 USD by 1000 to preserve precision up to one tenth of a cent.

```
{ price: 9990, currency: "USD" }
```

The model assumes that for a given currency value:

- The scale factor is consistent for a currency; i.e. same scaling factor for a given currency.
- The scale factor is a constant and known property of the currency; i.e applications can determine the scale factor from the currency.

When using this model, applications must be consistent in performing the appropriate scaling of the values.

For use cases of this model, see [Use Cases for Exact Precision Model](#) (page 26).

## Arbitrary Precision

To model monetary data using the arbitrary precision model, store the value in two fields:

1. In one field, encode the exact monetary value as a non-numeric data type; e.g., BinData or a string.
2. In the second field, store a double-precision floating point approximation of the exact value.

The following example uses the arbitrary precision model to store 9.99 USD for the price and 0.25 USD for the fee:

```
{
  price: { display: "9.99", approx: 9.990000000000002, currency: "USD" },
  fee: { display: "0.25", approx: 0.2499999999999999, currency: "USD" }
}
```

With some care, applications can perform range and sort queries on the field with the numeric approximation. However, the use of the approximation field for the query and sort operations requires that applications perform client-side post-processing to decode the non-numeric representation of the exact value and then filter out the returned documents based on the exact monetary value.

For use cases of this model, see [Use Cases for Arbitrary Precision Model](#) (page 26).

## Model Time Data

### Overview

MongoDB [stores times in UTC](#) (page 40) by default, and will convert any local time representations into this form. Applications that must operate or report on some unmodified local time value may store the time zone alongside the UTC timestamp, and compute the original local time in their application logic.

### Example

In the MongoDB shell, you can store both the current date and the current client's offset from UTC.

```
var now = new Date();
db.data.save( { date: now,
               offset: now.getTimezoneOffset() } );
```

You can reconstruct the original local time by applying the saved offset:

```
var record = db.data.findOne();  
var localNow = new Date( record.date.getTime() - ( record.offset * 60000 ) );
```

## 4 Data Model Reference

**Documents** (page 28) MongoDB stores all data in documents, which are JSON-style data structures composed of field-and-value pairs.

**Database References** (page 31) Discusses manual references and DBRefs, which MongoDB can use to represent relationships between documents.

**GridFS Reference** (page 34) Convention for storing large files in a MongoDB Database.

**ObjectId** (page 36) A 12-byte BSON type that MongoDB uses as the default value for its documents' `_id` field if the `_id` field is not specified.

**BSON Types** (page 38) Outlines the unique *BSON* types used by MongoDB. See [BSONspec.org](http://bsonspec.org)<sup>12</sup> for the complete BSON specification.

### 4.1 Documents

MongoDB stores all data in documents, which are JSON-style data structures composed of field-and-value pairs:

```
{ "item": "pencil", "qty": 500, "type": "no.2" }
```

Most user-accessible data structures in MongoDB are documents, including:

- All database records.
- Query `selectors`, which define what records to select for read, update, and delete operations.
- Update `definitions`, which define what fields to modify during an update.
- Index `specifications`, which define what fields to index.
- Data output by MongoDB for reporting and configuration, such as the output of the `serverStatus` and the *replica set configuration document*.

#### Document Format

MongoDB stores documents on disk in the *BSON* serialization format. BSON is a binary representation of *JSON* documents, though it contains more data types than JSON. For the BSON spec, see [bsonspec.org](http://bsonspec.org)<sup>13</sup>. See also **BSON Types** (page 38).

The `mongo` JavaScript shell and the MongoDB language drivers translate between BSON and the language-specific document representation.

#### Document Structure

MongoDB documents are composed of field-and-value pairs and have the following structure:

---

<sup>12</sup><http://bsonspec.org/>

<sup>13</sup><http://bsonspec.org/>

```
{
    field1: value1,
    field2: value2,
    field3: value3,
    ...
    fieldN: valueN
}
```

The value of a field can be any of the BSON *data types* (page 38), including other documents, arrays, and arrays of documents. The following document contains values of varying types:

```
var mydoc = {
    _id: ObjectId("5099803df3f4948bd2f98391"),
    name: { first: "Alan", last: "Turing" },
    birth: new Date('Jun 23, 1912'),
    death: new Date('Jun 07, 1954'),
    contribs: [ "Turing machine", "Turing test", "Turingery" ],
    views : NumberLong(1250000)
}
```

The above fields have the following data types:

- `_id` holds an *ObjectId*.
- `name` holds an *embedded document* that contains the fields `first` and `last`.
- `birth` and `death` hold values of the *Date* type.
- `contribs` holds an *array of strings*.
- `views` holds a value of the *NumberLong* type.

## Field Names

Field names are strings.

*Documents* (page 28) have the following restrictions on field names:

- The field name `_id` is reserved for use as a primary key; its value must be unique in the collection, is immutable, and may be of any type other than an array.
- The field names **cannot** start with the dollar sign (\$) character.
- The field names **cannot** contain the dot (.) character.
- The field names **cannot** contain the null character.

BSON documents may have more than one field with the same name. Most MongoDB interfaces, however, represent MongoDB with a structure (e.g. a hash table) that does not support duplicate field names. If you need to manipulate documents that have more than one field with the same name, see the `driver` documentation for your driver.

Some documents created by internal MongoDB processes may have duplicate fields, but *no* MongoDB process will ever add duplicate fields to an existing user document.

## Field Value Limit

For indexed collections, the values for the indexed fields have a Maximum Index Key Length limit. See Maximum Index Key Length for details.

## Document Limitations

Documents have the following attributes:

### Document Size Limit

The maximum BSON document size is 16 megabytes.

The maximum document size helps ensure that a single document cannot use excessive amount of RAM or, during transmission, excessive amount of bandwidth. To store documents larger than the maximum size, MongoDB provides the GridFS API. See `mongofiles` and the documentation for your driver for more information about GridFS.

### Document Field Order

MongoDB preserves the order of the document fields following write operations *except* for the following cases:

- The `_id` field is always the first field in the document.
- Updates that include renaming of field names may result in the reordering of fields in the document.

Changed in version 2.6: Starting in version 2.6, MongoDB actively attempts to preserve the field order in a document. Before version 2.6, MongoDB did not actively preserve the order of the fields in a document.

### The `_id` Field

The `_id` field has the following behavior and constraints:

- By default, MongoDB creates a unique index on the `_id` field during the creation of a collection.
- The `_id` field is always the first field in the documents. If the server receives a document that does not have the `_id` field first, then the server will move the field to the beginning.
- The `_id` field may contain values of any *BSON data type* (page 38), other than an array.

**Warning:** To ensure functioning replication, do not store values that are of the BSON regular expression type in the `_id` field.

The following are common options for storing values for `_id`:

- Use an *ObjectId* (page 36).
- Use a natural unique identifier, if available. This saves space and avoids an additional index.
- Generate an auto-incrementing number. See <http://docs.mongodb.org/manual/tutorial/create-an-auto-incrementing-field/>
- Generate a UUID in your application code. For a more efficient storage of the UUID values in the collection and in the `_id` index, store the UUID as a value of the BSON `BinData` type.

Index keys that are of the `BinData` type are more efficiently stored in the index if:

- the binary subtype value is in the range of 0-7 or 128-135, and
- the length of the byte array is: 0, 1, 2, 3, 4, 5, 6, 7, 8, 10, 12, 14, 16, 20, 24, or 32.
- Use your driver's BSON UUID facility to generate UUIDs. Be aware that driver implementations may implement UUID serialization and deserialization logic differently, which may not be fully compatible with other drivers. See your driver documentation<sup>14</sup> for information concerning UUID interoperability.

<sup>14</sup><http://api.mongodb.org/>

---

**Note:** Most MongoDB driver clients will include the `_id` field and generate an `ObjectId` before sending the insert operation to MongoDB; however, if the client sends a document without an `_id` field, the `mongod` will add the `_id` field and generate the `ObjectId`.

---

## Dot Notation

MongoDB uses the *dot notation* to access the elements of an array and to access the fields of an embedded document.

To access an element of an array by the zero-based index position, concatenate the array name with the dot (.) and zero-based index position, and enclose in quotes:

```
'<array>.<index>'
```

See also \$ positional operator for update operations and \$ projection operator when array index position is unknown.

To access a field of an embedded document with *dot-notation*, concatenate the embedded document name with the dot (.) and the field name, and enclose in quotes:

```
'<embedded document>.<field>'
```

### See also:

- *read-operations-embedded-documents* for dot notation examples with embedded documents.
- *read-operations-arrays* for dot notation examples with arrays.

## Additional Resources

- Thinking in Documents Part 1 (Blog Post)<sup>15</sup>

## 4.2 Database References

MongoDB does not support joins. In MongoDB some data is *denormalized*, or stored with related data in *documents* to remove the need for joins. However, in some cases it makes sense to store related information in separate documents, typically in different collections or databases.

MongoDB applications use one of two methods for relating documents:

- *Manual references* (page 32) where you save the `_id` field of one document in another document as a reference. Then your application can run a second query to return the related data. These references are simple and sufficient for most use cases.
- *DBRefs* (page 32) are references from one document to another using the value of the first document's `_id` field, collection name, and, optionally, its database name. By including these names, DBRefs allow documents located in multiple collections to be more easily linked with documents from a single collection.

To resolve DBRefs, your application must perform additional queries to return the referenced documents. Many drivers have helper methods that form the query for the DBRef automatically. The drivers <sup>16</sup> do not automatically resolve DBRefs into documents.

DBRefs provide a common format and type to represent relationships among documents. The DBRef format also provides common semantics for representing links between documents if your database must interact with multiple frameworks and tools.

---

<sup>15</sup><https://www.mongodb.com/blog/post/thinking-documents-part-1?jmp=docs>

<sup>16</sup> Some community supported drivers may have alternate behavior and may resolve a DBRef into a document automatically.

Unless you have a compelling reason to use DBRefs, use manual references instead.

## Manual References

### Background

Using manual references is the practice of including one *document's* `_id` field in another document. The application can then issue a second query to resolve the referenced fields as needed.

### Process

Consider the following operation to insert two documents, using the `_id` field of the first document as a reference in the second document:

```
original_id = ObjectId()

db.places.insert({
    "_id": original_id,
    "name": "Broadway Center",
    "url": "bc.example.net"
})

db.people.insert({
    "name": "Erin",
    "places_id": original_id,
    "url": "bc.example.net/Erin"
})
```

Then, when a query returns the document from the `people` collection you can, if needed, make a second query for the document referenced by the `places_id` field in the `places` collection.

### Use

For nearly every case where you want to store a relationship between two documents, use [manual references](#) (page 32). The references are simple to create and your application can resolve references as needed.

The only limitation of manual linking is that these references do not convey the database and collection names. If you have documents in a single collection that relate to documents in more than one collection, you may need to consider using DBRefs.

## DBRefs

### Background

DBRefs are a convention for representing a *document*, rather than a specific reference type. They include the name of the collection, and in some cases the database name, in addition to the value from the `_id` field.

### Format

DBRefs have the following fields:

### **\$ref**

The `$ref` field holds the name of the collection where the referenced document resides.

### **\$id**

The `$id` field contains the value of the `_id` field in the referenced document.

### **\$db**

*Optional.*

Contains the name of the database where the referenced document resides.

Only some drivers support `$db` references.

## **Example**

DBRef documents resemble the following document:

```
{ "$ref" : <value>, "$id" : <value>, "$db" : <value> }
```

Consider a document from a collection that stored a DBRef in a `creator` field:

```
{
  "_id" : ObjectId("5126bbf64aed4daf9e2ab771"),
  // .. application fields
  "creator" : {
    "$ref" : "creators",
    "$id" : ObjectId("5126bc054aed4daf9e2ab772"),
    "$db" : "users"
  }
}
```

The DBRef in this example points to a document in the `creators` collection of the `users` database that has `ObjectId("5126bc054aed4daf9e2ab772")` in its `_id` field.

**Note:** The order of fields in the DBRef matters, and you must use the above sequence when using a DBRef.

## **Driver Support for DBRefs**

**C** The C driver contains no support for DBRefs. You can traverse references manually.

**C++** The C++ driver contains no support for DBRefs. You can traverse references manually.

**C#** The C# driver supports DBRefs using the `MongoDBRef17` class and the `FetchDBRef18` method.

**Haskell** The Haskell driver contains no support for DBRefs. You can traverse references manually.

**Java** The `DBRef19` class provides support for DBRefs from Java.

**JavaScript** The mongo shell's JavaScript interface provides a DBRef.

**Node.js** The Node.js driver supports DBRefs using the `DBRef20` class and the `dereference21` method.

**Perl** The Perl driver contains no support for DBRefs. You can traverse references manually or use the `MongoDBx::AutoDeref22` CPAN module.

<sup>17</sup><http://api.mongodb.org/csharp/current/html/46c356d3-ed06-a6f8-42fa-e0909ab64ce2.htm>

<sup>18</sup><http://api.mongodb.org/csharp/current/html/1b0b8f48-ba98-1367-0a7d-6e01c8df436f.htm>

<sup>19</sup><http://api.mongodb.org/java/current/com/mongodb/DBRef.html>

<sup>20</sup>[http://mongodb.github.io/node-mongodb-native/api-bson-generated/db\\_ref.html](http://mongodb.github.io/node-mongodb-native/api-bson-generated/db_ref.html)

<sup>21</sup><http://mongodb.github.io/node-mongodb-native/api-generated/db.html#dereference>

<sup>22</sup><http://search.cpan.org/dist/MongoDBx-AutoDeref/>

**PHP** The PHP driver supports DBRefs, including the optional `$db` reference, using the `MongoDBRef`<sup>23</sup> class.

**Python** The Python driver supports DBRefs using the `DBRef`<sup>24</sup> class and the `dereference`<sup>25</sup> method.

**Ruby** The Ruby driver supports DBRefs using the `DBRef`<sup>26</sup> class and the `dereference`<sup>27</sup> method.

**Scala** The Scala driver contains no support for DBRefs. You can traverse references manually.

## Use

In most cases you should use the *manual reference* (page 32) method for connecting two or more related documents. However, if you need to reference documents from multiple collections, consider using DBRefs.

## 4.3 GridFS Reference

*GridFS* stores files in two collections:

- `chunks` stores the binary chunks. For details, see *The chunks Collection* (page 34).
- `files` stores the file's metadata. For details, see *The files Collection* (page 35).

GridFS places the collections in a common bucket by prefixing each with the bucket name. By default, GridFS uses two collections with names prefixed by `fs` bucket:

- `fs.files`
- `fs.chunks`

You can choose a different bucket name than `fs`, and create multiple buckets in a single database.

**See also:**

*GridFS* (page 10) for more information about GridFS.

### The chunks Collection

Each document in the `chunks` collection represents a distinct chunk of a file as represented in the *GridFS* store. The following is a prototype document from the `chunks` collection.:

```
{  
  "_id" : <ObjectId>,  
  "files_id" : <ObjectId>,  
  "n" : <num>,  
  "data" : <binary>  
}
```

A document from the `chunks` collection contains the following fields:

`chunks._id`

The unique *ObjectId* of the chunk.

`chunks.files_id`

The `_id` of the “parent” document, as specified in the `files` collection.

<sup>23</sup><http://www.php.net/manual/en/class.mongodb.php/>

<sup>24</sup><http://api.mongodb.org/python/current/api/bson/dbref.html>

<sup>25</sup><http://api.mongodb.org/python/current/api/pymongo/database.html#pymongo.database.Database.dereference>

<sup>26</sup><http://api.mongodb.org/ruby/current/BSON/DBRef.html>

<sup>27</sup>[http://api.mongodb.org/ruby/current/Mongo/DB.html#dereference-instance\\_method](http://api.mongodb.org/ruby/current/Mongo/DB.html#dereference-instance_method)

### `chunks.n`

The sequence number of the chunk. GridFS numbers all chunks, starting with 0.

### `chunks.data`

The chunk's payload as a *BSON* binary type.

The `chunks` collection uses a *compound index* on `files_id` and `n`, as described in [GridFS Index](#) (page 11).

## The `files` Collection

Each document in the `files` collection represents a file in the *GridFS* store. Consider the following prototype of a document in the `files` collection:

```
{  
  "_id" : <ObjectId>,  
  "length" : <num>,  
  "chunkSize" : <num>,  
  "uploadDate" : <timestamp>,  
  "md5" : <hash>,  
  
  "filename" : <string>,  
  "contentType" : <string>,  
  "aliases" : <string array>,  
  "metadata" : <dataObject>,  
}
```

Documents in the `files` collection contain some or all of the following fields. Applications may create additional arbitrary fields:

### `files._id`

The unique ID for this document. The `_id` is of the data type you chose for the original document. The default type for MongoDB documents is *BSON ObjectId*.

### `files.length`

The size of the document in bytes.

### `files.chunkSize`

The size of each chunk. GridFS divides the document into chunks of the size specified here. The default size is 255 kilobytes.

Changed in version 2.4.10: The default chunk size changed from 256k to 255k.

### `files.uploadDate`

The date the document was first stored by GridFS. This value has the `Date` type.

### `files.md5`

An MD5 hash returned by the `filemd5` command. This value has the `String` type.

### `files.filename`

Optional. A human-readable name for the document.

### `files.contentType`

Optional. A valid MIME type for the document.

### `files.aliases`

Optional. An array of alias strings.

### `files.metadata`

Optional. Any additional information you want to store.

## 4.4 ObjectId

### Overview

*ObjectId* is a 12-byte *BSON* type, constructed using:

- a 4-byte value representing the seconds since the Unix epoch,
- a 3-byte machine identifier,
- a 2-byte process id, and
- a 3-byte counter, starting with a random value.

In MongoDB, documents stored in a collection require a unique `_id` field that acts as a *primary key*. Because ObjectIds are small, most likely unique, and fast to generate, MongoDB uses ObjectIds as the default value for the `_id` field if the `_id` field is not specified. MongoDB clients should add an `_id` field with a unique ObjectId. However, if a client does not add an `_id` field, `mongod` will add an `_id` field that holds an ObjectId.

Using ObjectIds for the `_id` field provides the following additional benefits:

- in the `mongo` shell, you can access the creation time of the `ObjectId`, using the `getTimestamp()` method.
- sorting on an `_id` field that stores `ObjectId` values is roughly equivalent to sorting by creation time.

---

**Important:** The relationship between the order of `ObjectId` values and generation time is not strict within a single second. If multiple systems, or multiple processes or threads on a single system generate values, within a single second; `ObjectId` values do not represent a strict insertion order. Clock skew between clients can also result in non-strict ordering even for values, because client drivers generate `ObjectId` values, *not* the `mongod` process.

---

Also consider the [Documents](#) (page 28) section for related information on MongoDB's document orientation.

### ObjectId()

The `mongo` shell provides the `ObjectId()` wrapper class to generate a new ObjectId, and to provide the following helper attribute and methods:

- `str`  
The hexadecimal string representation of the object.
  - `getTimestamp()`  
Returns the timestamp portion of the object as a Date.
  - `toString()`  
Returns the JavaScript representation in the form of a string literal “`ObjectId(...)`”.
- Changed in version 2.2: In previous versions `toString()` returns the hexadecimal string representation, which as of version 2.2 can be retrieved by the `str` property.
- `valueOf()`  
Returns the representation of the object as a hexadecimal string. The returned string is the `str` attribute.
- Changed in version 2.2: In previous versions, `valueOf()` returns the object.

## Examples

Consider the following uses `ObjectId()` class in the mongo shell:

### Generate a new ObjectId

To generate a new ObjectId, use the `ObjectId()` constructor with no argument:

```
x = ObjectId()
```

In this example, the value of `x` would be:

```
ObjectId("507f1f77bcf86cd799439011")
```

To generate a new ObjectId using the `ObjectId()` constructor with a unique hexadecimal string:

```
y = ObjectId("507f191e810c19729de860ea")
```

In this example, the value of `y` would be:

```
ObjectId("507f191e810c19729de860ea")
```

- To return the timestamp of an `ObjectId()` object, use the `getTimestamp()` method as follows:

### Convert an ObjectId into a Timestamp

To return the timestamp of an `ObjectId()` object, use the `getTimestamp()` method as follows:

```
ObjectId("507f191e810c19729de860ea").getTimestamp()
```

This operation will return the following Date object:

```
ISODate("2012-10-17T20:46:22Z")
```

### Convert ObjectIds into Strings

Access the `str` attribute of an `ObjectId()` object, as follows:

```
ObjectId("507f191e810c19729de860ea").str
```

This operation will return the following hexadecimal string:

```
507f191e810c19729de860ea
```

To return the hexadecimal string representation of an `ObjectId()`, use the `valueOf()` method as follows:

```
ObjectId("507f191e810c19729de860ea").valueOf()
```

This operation returns the following output:

```
507f191e810c19729de860ea
```

To return the string representation of an `ObjectId()` object, use the `toString()` method as follows:

```
ObjectId("507f191e810c19729de860ea").toString()
```

This operation will return the following output:

```
ObjectId("507f191e810c19729de860ea")
```

## 4.5 BSON Types

*BSON* is a binary serialization format used to store documents and make remote procedure calls in MongoDB. The *BSON* specification is located at [bsonspec.org](http://bsonspec.org)<sup>28</sup>.

*BSON* supports the following data types as values in documents. Each data type has a corresponding number that can be used with the `$type` operator to query documents by *BSON* type.

Type	Number	Notes
Double	1	
String	2	
Object	3	
Array	4	
Binary data	5	
Undefined	6	Deprecated.
Object id	7	
Boolean	8	
Date	9	
Null	10	
Regular Expression	11	
JavaScript	13	
Symbol	14	
JavaScript (with scope)	15	
32-bit integer	16	
Timestamp	17	
64-bit integer	18	
Min key	255	Query with <code>-1</code> .
Max key	127	

To determine a field's type, see *check-types-in-shell*.

If you convert *BSON* to JSON, see the [Extended JSON reference](#).

## Comparison/Sort Order

When comparing values of different *BSON* types, MongoDB uses the following comparison order, from lowest to highest:

1. MinKey (internal type)
2. Null
3. Numbers (ints, longs, doubles)
4. Symbol, String
5. Object
6. Array
7. BinData
8. ObjectId
9. Boolean

---

<sup>28</sup><http://bsonspec.org/>

10. Date
11. Timestamp
12. Regular Expression
13. MaxKey (internal type)

MongoDB treats some types as equivalent for comparison purposes. For instance, numeric types undergo conversion before comparison.

Changed in version 3.0.0: Date objects sort before Timestamp objects. Previously Date and Timestamp objects sorted together.

The comparison treats a non-existent field as it would an empty BSON Object. As such, a sort on the `a` field in documents `{ }` and `{ a: null }` would treat the documents as equivalent in sort order.

With arrays, a less-than comparison or an ascending sort compares the smallest element of arrays, and a greater-than comparison or a descending sort compares the largest element of the arrays. As such, when comparing a field whose value is a single-element array (e.g. `[ 1 ]`) with non-array fields (e.g. `2`), the comparison is between `1` and `2`. A comparison of an empty array (e.g. `[ ]`) treats the empty array as less than `null` or a missing field.

MongoDB sorts `BinData` in the following order:

1. First, the length or size of the data.
2. Then, by the BSON one-byte subtype.
3. Finally, by the data, performing a byte-by-byte comparison.

The following sections describe special considerations for particular BSON types.

## ObjectId

ObjectIds are: small, likely unique, fast to generate, and ordered. These values consists of 12-bytes, where the first four bytes are a timestamp that reflect the ObjectId's creation. Refer to the [ObjectId](#) (page 36) documentation for more information.

## String

BSON strings are UTF-8. In general, drivers for each programming language convert from the language's string format to UTF-8 when serializing and deserializing BSON. This makes it possible to store most international characters in BSON strings with ease.<sup>29</sup> In addition, MongoDB `$regex` queries support UTF-8 in the regex string.

## Timestamps

BSON has a special timestamp type for *internal* MongoDB use and is **not** associated with the regular [Date](#) (page 40) type. Timestamp values are a 64 bit value where:

- the first 32 bits are a `time_t` value (seconds since the Unix epoch)
- the second 32 bits are an incrementing `ordinal` for operations within a given second.

Within a single `mongod` instance, timestamp values are always unique.

In replication, the `oplog` has a `ts` field. The values in this field reflect the operation time, which uses a BSON timestamp value.

---

<sup>29</sup> Given strings using UTF-8 character sets, using `sort()` on strings will be reasonably correct. However, because internally `sort()` uses the C++ `strcmp` api, the sort order may handle some characters incorrectly.

---

**Note:** The BSON timestamp type is for *internal* MongoDB use. For most cases, in application development, you will want to use the BSON date type. See [Date](#) (page 40) for more information.

---

If you insert a document containing an empty BSON timestamp in a top-level field, the MongoDB server will replace that empty timestamp with the current timestamp value. For example, if you create an insert a document with a timestamp value, as in the following operation:

```
var a = new Timestamp();  
  
db.test.insert( { ts: a } );
```

Then, the `db.test.find()` operation will return a document that resembles the following:

```
{ "_id" : ObjectId("542c2b97bac0595474108b48"), "ts" : Timestamp(1412180887, 1) }
```

If `ts` were a field in an embedded document, the server would have left it as an empty timestamp value.

Changed in version 2.6: Previously, the server would only replace empty timestamp values in the first two fields, including `_id`, of an inserted document. Now MongoDB will replace any top-level field.

## Date

BSON Date is a 64-bit integer that represents the number of milliseconds since the Unix epoch (Jan 1, 1970). This results in a representable date range of about 290 million years into the past and future.

The [official BSON specification](#)<sup>30</sup> refers to the BSON Date type as the *UTC datetime*.

Changed in version 2.0: BSON Date type is signed.<sup>31</sup> Negative values represent dates before 1970.

---

### Example

Construct a Date using the `new Date()` constructor in the mongo shell:

```
var mydate1 = new Date()
```

---

### Example

Construct a Date using the `ISODATE()` constructor in the mongo shell:

```
var mydate2 = ISODATE()
```

---

### Example

Return the Date value as string:

```
mydate1.toString()
```

---

### Example

Return the month portion of the Date value; months are zero-indexed, so that January is month 0:

---

<sup>30</sup><http://bsonspec.org/#/specification>

<sup>31</sup> Prior to version 2.0, Date values were incorrectly interpreted as *unsigned* integers, which affected sorts, range queries, and indexes on Date fields. Because indexes are not recreated when upgrading, please re-index if you created an index on Date values with an earlier version, and dates before 1970 are relevant to your application.

mydate1.getMonth()

---